

Deliverable 2 — Sprint Reflection

Customizable Subscription Box Scheduler • Problem Statement #35 | Retail, E-Commerce & Finance

Name	SRN	Course	Jira Project
Prajwal G Devadiga	PES1UG24CS590	Software Engineering - Lab 2	SBOX

Sprint 1 committed 18 points across 3 stories; Sprint 2 committed 16 points across 3 stories. EPIC-1 totals 34 points.

1. Did your estimations reflect the actual effort?

No story was actually implemented, so there is no real effort figure to compare against - the honest test is whether the estimates were consistent relative to each other. On the large items they were: US-1.2 and US-1.5 both drew a 5 / 8 / 13 spread and settled at 8, with the high estimates coming from genuine uncertainty (re-picking a box whose stock is already reserved; bounded retries across a suspension state) rather than from the amount of code.

The clear error is US-1.6, a unanimous 3 that nobody questioned. NFR-001 requires 10,000 shipping labels exported in under 60 seconds, which turns a list-generation story into a throughput problem - it belongs at 8. The estimate was made against the story title rather than the requirement behind it.

2. Was your backlog well-prioritized?

Yes for delivery order. The sequence follows the real dependency chain - preference tags have to exist before a box can be recommended, a box has to be customisable before a cut-off means anything, and the box has to be locked before it can be billed or picked - so US-1.1 to US-1.3 fill Sprint 1 and the three stories that consume their output fill Sprint 2.

The weakness is the priority field itself: four of the six stories are High, so the label carries little information and the sprint assignment does the real prioritising. US-1.3 is under-rated at Medium, since it is the only story stopping the warehouse pick list from moving after the manifest has been cut.

3. How did your simulated sprint align with your plan?

Exactly, and that is the least interesting thing about it. Sprint 1 committed 18 points across 3 stories and Sprint 2 committed 16 points across 3; everything moved To Do -> In Progress -> Done inside the session, so both closed at 100% of commitment with nothing carried over.

That is an artefact of the simulation, not good planning. No code was written, so nothing could block and no estimate could be proved wrong - US-1.6 would have been the one to slip. Two sprints closing at exactly 100% would normally read as a team under-committing.

4. What insights did the burndown chart give about your team's capacity?

The remaining-work line does not track the guideline at all: flat while the stories sit in In Progress, then near-vertical as they are all closed at the end. In a real sprint that shape signals work being integrated in one late batch rather than finished incrementally; here it is simply the effect of compressing a one-week sprint into a lab session.

On capacity, 18 then 16 points averages 17 - but two data points is far too small to plan from, and the two sprints were sized so alike that the difference between them says nothing about the team either.